

Homework2

Justin Lau

10/28/2020

Libraries

```
library(tidyverse)
```

```
## — Attaching packages — tidyverse 1.3.0 —
```

```
## ✓ ggplot2 3.3.2    ✓ purrr 0.3.4
## ✓ tibble 3.0.4     ✓ dplyr 1.0.2
## ✓ tidyr 1.1.2      ✓ stringr 1.4.0
## ✓ readr 1.4.0      ✓ forcats 0.5.0
```

```
## — Conflicts — tidyverse_conflicts() —
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
```

```
library(reshape2)
```

```
##
## Attaching package: 'reshape2'
```

```
## The following object is masked from 'package:tidyr':
##
## smiths
```

```
library(tree)
```

```
## Registered S3 method overwritten by 'tree':
## method      from
## print.tree cli
```

```
library(plyr)
```

```
## -----
```

```
## You have loaded plyr after dplyr - this is likely to cause problems.
## If you need functions from both plyr and dplyr, please load plyr first, then dplyr:
## library(plyr); library(dplyr)
```

```
## -----
```

```
##  
## Attaching package: 'plyr'
```

```
## The following objects are masked from 'package:dplyr':  
##  
##   arrange, count, desc, failwith, id, mutate, rename, summarise,  
##   summarize
```

```
## The following object is masked from 'package:purrr':  
##  
##   compact
```

```
library(class)  
library(rpart)  
library(maptree)
```

```
## Loading required package: cluster
```

```
library(ROCR)
```

```
spam <- read_table2("spambase.tab", guess_max=2000)
```

```
##  
## — Column specification —————  
## cols(  
##   .default = col_double()  
## )  
## i Use `spec()` for the full column specifications.
```

```
spam <- spam %>%  
  mutate(y = factor(y, levels=c(0,1), labels=c("good", "spam"))) %>%  
  mutate_at(.vars=vars(-y), .funs=scale) # scale others
```

```
calc_error_rate <- function(predicted.value, true.value){  
  return(mean(true.value!=predicted.value))  
}
```

```
records = matrix(NA, nrow=3, ncol=2)  
colnames(records) <- c("train.error", "test.error")  
rownames(records) <- c("knn", "tree", "logistic")
```

```
set.seed(1)
test.indices = sample(1:nrow(spam), 1000)
spam.train=spam[-test.indices,]
spam.test=spam[test.indices,]
```

```
nfold = 10
set.seed(1)
folds = seq.int(nrow(spam.train)) %>% ## sequential obs ids
cut(breaks = nfold, labels=FALSE) %>% ## sequential fold ids
sample ## random fold ids
```

```
do.chunk <- function(chunkid, folddef, Xdat, Ydat, k){
  train = (folddef!=chunkid)
  Xtr = Xdat[train,]
  Ytr = Ydat[train]
  Xvl = Xdat[!train,]
  Yvl = Ydat[!train]
  ## get classifications for current training chunks
  predYtr = knn(train = Xtr, test = Xtr, cl = Ytr, k = k)
  ## get classifications for current test chunk
  predYvl = knn(train = Xtr, test = Xvl, cl = Ytr, k = k)
  data.frame(fold = chunkid, train.error = calc_error_rate(predYtr, Ytr),
    val.error = calc_error_rate(predYvl, Yvl))
}
```

Question 1

```
error.folds <- NULL
YTrain <- spam.train$y
XTrain <- spam.train %>% select(-y)

YTest <- spam.test$y
XTest <- spam.test %>% select(-y)

set.seed(1)
kvec = c(1, seq(10, 50, length.out=5))
for (j in kvec){
  tmp <- ldply(1:nfold, do.chunk, folddef = folds, Xdat = XTrain, Ydat = YTrain, k = j)
  tmp$neighbors <- j
  error.folds <- rbind(error.folds,tmp)
}
error.folds
```

##	fold	train.error	val.error	neighbors
## 1	1	0.0006172840	0.11080332	1
## 2	2	0.0000000000	0.11944444	1
## 3	3	0.0006170935	0.08055556	1
## 4	4	0.0000000000	0.08055556	1
## 5	5	0.0006170935	0.10833333	1
## 6	6	0.0006170935	0.11111111	1
## 7	7	0.0003085467	0.07777778	1
## 8	8	0.0000000000	0.11666667	1
## 9	9	0.0003085467	0.10000000	1
## 10	10	0.0003085467	0.13055556	1
## 11	1	0.0824074074	0.08864266	10
## 12	2	0.0823819809	0.11111111	10
## 13	3	0.0805307004	0.08888889	10
## 14	4	0.0774452330	0.10000000	10
## 15	5	0.0755939525	0.09722222	10
## 16	6	0.0762110460	0.10277778	10
## 17	7	0.0805307004	0.05833333	10
## 18	8	0.0789879667	0.09444444	10
## 19	9	0.0759024992	0.11111111	10
## 20	10	0.0786794199	0.11388889	10
## 21	1	0.0919753086	0.09418283	20
## 22	2	0.0944153039	0.11944444	20
## 23	3	0.0956494909	0.08055556	20
## 24	4	0.0934896637	0.08888889	20
## 25	5	0.0888614625	0.12500000	20
## 26	6	0.0882443690	0.11111111	20
## 27	7	0.0965751311	0.06944444	20
## 28	8	0.0907127430	0.10555556	20
## 29	9	0.0931811169	0.12777778	20
## 30	10	0.0910212897	0.10000000	20
## 31	1	0.0993827160	0.10249307	30
## 32	2	0.1024375193	0.12500000	30
## 33	3	0.1052144400	0.10000000	30
## 34	4	0.1030546128	0.10555556	30
## 35	5	0.0993520518	0.11666667	30
## 36	6	0.0984264116	0.10833333	30
## 37	7	0.1033631595	0.07777778	30
## 38	8	0.0971922246	0.12777778	30
## 39	9	0.1012033323	0.11944444	30
## 40	10	0.0990435051	0.10833333	30
## 41	1	0.1055555556	0.11357341	40
## 42	2	0.1058315335	0.11666667	40
## 43	3	0.1104597346	0.11111111	40
## 44	4	0.1052144400	0.10555556	40
## 45	5	0.1101511879	0.12222222	40
## 46	6	0.1073742672	0.12222222	40
## 47	7	0.1098426412	0.08055556	40
## 48	8	0.1021289725	0.13055556	40
## 49	9	0.1098426412	0.12500000	40
## 50	10	0.1033631595	0.09722222	40
## 51	1	0.1111111111	0.11080332	50
## 52	2	0.1129281086	0.12222222	50

```
## 53      3 0.1135452021 0.11388889      50
## 54      4 0.1110768281 0.11111111      50
## 55      5 0.1104597346 0.12222222      50
## 56      6 0.1116939216 0.11944444      50
## 57      7 0.1141622956 0.08055556      50
## 58      8 0.1110768281 0.14722222      50
## 59      9 0.1082999074 0.12222222      50
## 60     10 0.1089170009 0.10555556      50
```

```
errors = melt(error.folds, id.vars=c('fold', 'neighbors'), value.name='error')
val.error.means = errors %>%
  filter(variable=='val.error') %>%
  group_by(neighbors, variable) %>%
  summarise_each(funs(mean), error) %>%
  ungroup() %>%
  filter(error==min(error))
```

```
## Warning: `summarise_each()` is deprecated as of dplyr 0.7.0.
## Please use `across()` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_warnings()` to see where this warning was generated.
```

```
## Warning: `funs()` is deprecated as of dplyr 0.8.0.
## Please use a list of either functions or lambdas:
##
##   # Simple named list:
##   list(mean = mean, median = median)
##
##   # Auto named with `tibble::lst()`:
##   tibble::lst(mean, median)
##
##   # Using lambdas
##   list(~ mean(., trim = .2), ~ median(., na.rm = TRUE))
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_warnings()` to see where this warning was generated.
```

```
numneighbor = max(val.error.means$neighbors)
numneighbor
```

```
## [1] 10
```

```
sum(is.na(spam.test))
```

```
## [1] 0
```

Question 2

```

train_error <- do.chunk(6, folds, Xdat = XTrain, Ydat = YTrain, k = 10)

pred.YTest = knn(train=XTrain, test=XTest, cl=YTrain, k=numneighbor)
test_error <- calc_error_rate(pred.YTest,YTest)

records <- replace(records,1, train_error$train.error)
records <- replace(records,4, test_error)

```

Question 3

```

controls <- tree.control(nobs= nrow(spam.train),mincut=5, mindev=1e-5)
spamtree <- tree(y~., spam.train, control=controls)
summary(spamtree)

```

```

##
## Classification tree:
## tree(formula = y ~ ., data = spam.train, control = controls)
## Variables actually used in tree construction:
## [1] "char_freq..4"          "word_freq_remove"
## [3] "char_freq..3"          "word_freq_free"
## [5] "word_freq_george"      "word_freq_hp"
## [7] "capital_run_length_longest" "word_freq_receive"
## [9] "capital_run_length_average" "word_freq_credit"
## [11] "word_freq_your"        "word_freq_mail"
## [13] "word_freq_re"          "word_freq_you"
## [15] "capital_run_length_total" "word_freq_will"
## [17] "word_freq_edu"         "word_freq_people"
## [19] "word_freq_money"       "word_freq_our"
## [21] "word_freq_1999"        "word_freq_make"
## [23] "char_freq.."           "word_freq_data"
## [25] "word_freq_all"         "word_freq_over"
## [27] "char_freq..1"          "word_freq_project"
## [29] "word_freq_meeting"     "word_freq_internet"
## [31] "word_freq_650"         "word_freq_hpl"
## [33] "char_freq..5"          "word_freq_email"
## [35] "word_freq_business"    "word_freq_order"
## [37] "word_freq_address"
## Number of terminal nodes: 129
## Residual mean deviance: 0.1071 = 371.7 / 3472
## Misclassification error rate: 0.02555 = 92 / 3601

```

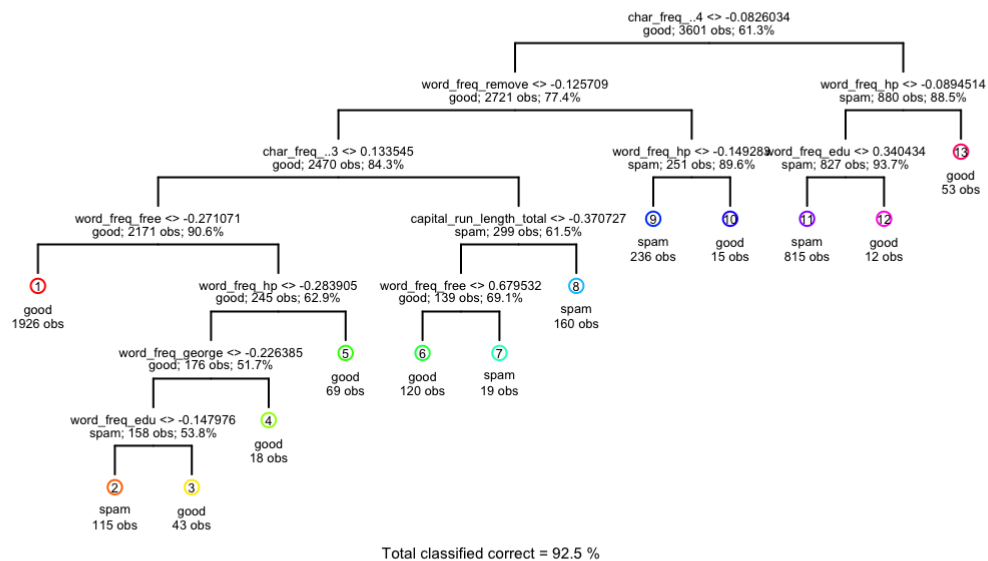
There was 133 missclassified observations and there are 141 leaf nodes.

Question 4

```

prune <- prune.tree(spamtree,best=10, method = 'misclass')
draw.tree(prune, nodeinfo = TRUE, cex = 0.4)


```



Question 5

```
cv <- cv.tree(spamtree, rand=folds, FUN = prune.misclass, K=10)

plot(cv$size, cv$dev, xlab = "Tree size", ylab = "Cross Validation Error")
```



```
best.size.cv = min(cv$size[cv$dev == 351])
```

```
## Warning in min(cv$size[cv$dev == 351]): no non-missing arguments to min;  
## returning Inf
```

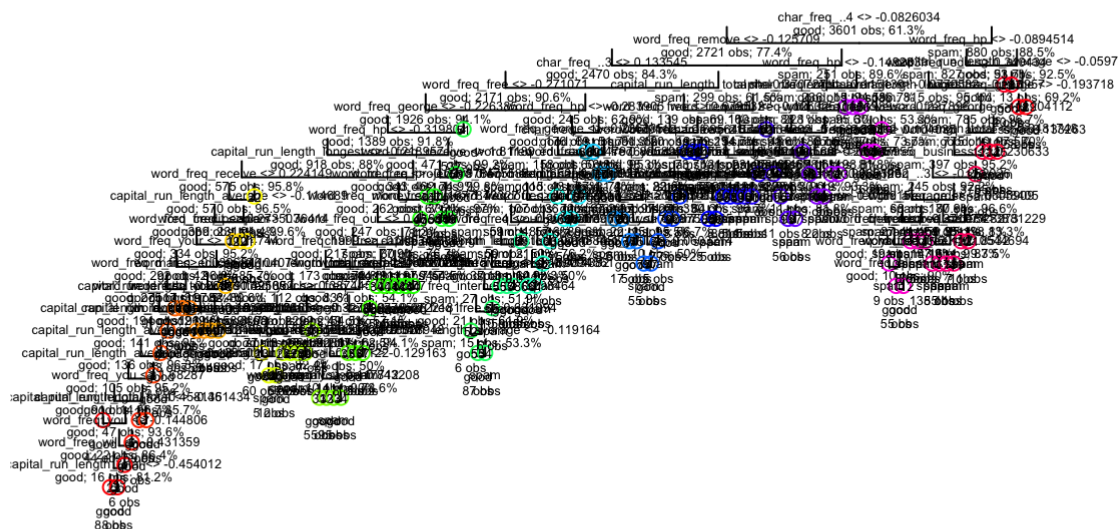
```
best.size.cv
```

```
## [1] Inf
```

```
spamtrees.pruned <- prune.misclass(spamtrees, best = best.size.cv)
```

```
## Warning in prune.tree(tree = spamtrees, best = best.size.cv, method =  
## "misclass"): best is bigger than tree size
```

```
draw.tree(spamtrees.pruned, nodeinfo = TRUE, cex = 0.4)
```

35 is

the optimal amount tree size.

Question 6

```
predict.pruned.test <- predict(spamtree.pruned, spam.test, type = 'class')
predict.pruned.train <- predict(spamtree.pruned, spam.train, type = 'class')
prune.test.error <- calc_error_rate(predict.pruned.test, spam.test$y)
prune.train.error <- calc_error_rate(predict.pruned.train, spam.train$y)

records <- replace(records,2, prune.train.error)
records <- replace(records,5, prune.test.error)
```

Question 7 will be written out in a separate segment

Question 8

```
glm.fit <- glm(y~., data=spam.train, family = binomial)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
prob.train = predict(glm.fit, spam.train, type="response")

prob.test = predict(glm.fit, spam.test, type="response")

spam.test <-
  spam.test %>%
  mutate(predSPAM = as.factor(ifelse(prob.test <= .5, "good", "spam")))

spam.train <-
  spam.train %>%
  mutate(predSPAM = as.factor(ifelse(prob.train <= .5, "good", "spam")))
```

```
log.test.error <- calc_error_rate(spam.test$predSPAM, spam.test$y)
log.train.error <- calc_error_rate(spam.train$predSPAM, spam.train$y)


records <- replace(records,3, log.train.error)
records <- replace(records,6, log.test.error)
print(records)
```

```
##          train.error test.error
## knn          0.08083925      0.103
## tree          0.02471536      0.098
## logistic      0.06803666      0.086
```

Question 9

```
pred <- prediction(prob.train, spam.train$y)
perf = performance(pred, measure="tpr", x.measure="fpr")
plot(perf, col=2, lwd=3, main="ROC curve")
abline(0,1)
```

ROC curve



```
auc = performance(pred, "auc")@y.values
auc
```

```
## [[1]]
## [1] 0.9772759
```

We are more worried with false positive rates that are too large as that would filter out emails that could be potentially important to a client. While having a large true positive rate that is too small would mean that not a lot of spam is being filtered out, almost making the filter worthless, having a large false positive rate would be more of a detriment more than anything else.